(a) Leave-one-slide-out partition of the 193 labelled cores

	slide001 36 cores	slide002 29 cores	slide003 28 cores	slide005 31 cores	slide006 44 cores	slide007 25 cores
Fold 0	Test	Val	Train	Train	Train	Train
Fold 1	Train	Test	Val	Train	Train	Train
Fold 2	Train	Train	Test	Val	Train	Train
Fold 3	Train	Train	Train	Test	Val	Train
Fold 4	Train	Train	Train	Train	Test	Val
Fold 5	Val	Train	Train	Train	Train	Test

Every core is tested exactly once (pooled n = 193).
A test slide is scored once, after training and checkpoint selection.

(b) Inside every fold

Training pool: the four remaining slides (118 to 136 cores)

↓

35 cores (fixed seed), 35 tumour-biased 512 × 512 patches each: 1,225 patches

↓

Validation slide: early stopping, learning-rate schedule, checkpoint choice

↓

Test slide: 20 random patches per core, scored once (Dice, IoU, HD95, ASD)

Same folds, patches and scoring for every model: three FrozenSE encoders, two U-Nets, nnU-Net and zero-shot MobileSAM